A Social Recommender System For Wines

Peter Chamberlin
Undergraduate Student
Department of Computer Science
Birkbeck, University of London

March 28, 2013

Abstract

This is the paper's abstract ...

Contents

1 Introduction

Recommender Systems.

Since their origin in the mid-1990s with systems such as Tapestry [6] and GroupLens [9], recommender systems have become ubiquitous on the World Wide Web, being employed by some of the worlds largest online businesses as core parts of their offering to users.

Companies such as Amazon, Netflix, Facebook and Twitter use recommender systems to make all manner of suggestions to their users. These recommendations include such things as products, movies, news stories and other interesting users.

It is the growth of the Web, which is now ubiquitous itself, that has given companies the ability to draw on unprecedented amounts of data about their users' preferences. At the same time the Web has made it easier than ever to reach their users with tailored suggestions.

Amazon's system of product recommendations using item-to-item collaborative filtering is regarded as a "killer feature"[7], and is one of the defining features of the Amazon brand experience. Amazon state their mission to be, "to delight our customers by allowing them to serendipitously discover great products"[7].

Netflix's movie recommender system *Cinematch* "Netflix Prize" competition

Social Recommender Systems.

In recent years social applications have dominated the Web. Networks like Facebook and Twitter have become massive global businesses as Internet users share more and more of their lives online. In such a context recommender systems are able to look beyond users' relationships to product and services, being able instead to interrogate qualitative data about the relationships between specific people. Facebook describe this as the "social graph".

CITATIONS NEEDED!

These systems, "a class of recommender systems that target the social media domain"[**?**], represent the current state of the art.

Recommending Wines.

Recommender systems for wines are not a new idea, being typical of the kind of item many systems are designed to recommend. Burke developed the VintageExchange FindMe recommender system in 1999[1], and there is at least one patent pending with the WIPO for a wine recommender system[**?**].

As a knowledge-based recommender, Burke's FindMe system, "required approximately one person-month of knowledge engineering effort"[2].

The Tetherless World Wine Agent (TWWA), by Patton and McGuinness[**?**]. The TWWA project is primarily concerned with knowledge representaion and the Semantic Web, presenting a common and collaborative ontology for wine with which users can share wine recommendations across their social networks[11]. The system does not automatically tailor recommendations to users, although this is stated as a target for future work[11].

Service Oriented Architecture.

I have chosen to implement my system as a service, such

Aims and Objectives.

I aim to produce a recommender system for wines which takes advantage of both ratings and tasting notes to suggest the most interesting wines and users. I intend the system to ignore wine attributes, such as grape variety and colour, preferring to make recommendations based on a hybrid of pure collaborative filtering and a content-based filtering approach using user-submitted and expert tasting notes.

Rather than implement a full graphical interface for the system I have chosen to develop an HTTP API.

In doing so I will explore the field of recommender systems,

Why are these systems interesting - Benefits - Challenges

Typical applications... - Movies - Products (i.e. Amazon)

Applications in wine domain - what's the same - what's different

What will this project do? - implement a recommender system for wines - exploration of techniques etc.

2 Literature Review

The term "recommender system" was coined by Resnick and Varian [10] to describe a system that "assists and augments" the "natural social process" of recommendation, preferring it to the more narrow term "collaborative filtering" used by Goldberg et al. [6] to describe their Tapestry system.

There are now several main categories of filtering technique employed in recommender systems. Burke [4] looks at five of these: collaborative, content-based, demographic, utiliy-based and knowledge-based (Burke, 2002).

Collaborative Filtering.

— off the web (?) — on the web...
– What are the methods employed in recommender systems?
– Collaborative Filtering - User-based filtering - Item-based filtering
– Content-Based Filtering - Variants
– Characteristics of the domain
— Cold start problem — Sparsity problem — ... etc.

3 Development Method

4 The Sommelier System

Data Cleanup

The data source I have used for my project is the wines database belonging to Decanter.com[**?**]. The database contains nearly 40,000 professional ratings and tasting notes for wines from as far back as 1986, featuring vintages as far back as 1917.

The original database is highly inconsistent, displaying a mixture of design approaches and a variable quality of data. This is consistent with the fact that the database has been developed over a long period of time by a number of different developers with varying levels of skill, and that wine journalists making entries into the database have taken a number of idiosyncratic approaches to data entry.

Nevertheless there is a great deal of useful and interesting information in the database, as it contains ratings and/or tastings for over 33,000 wines.

Figure 1: Decanter.com Database

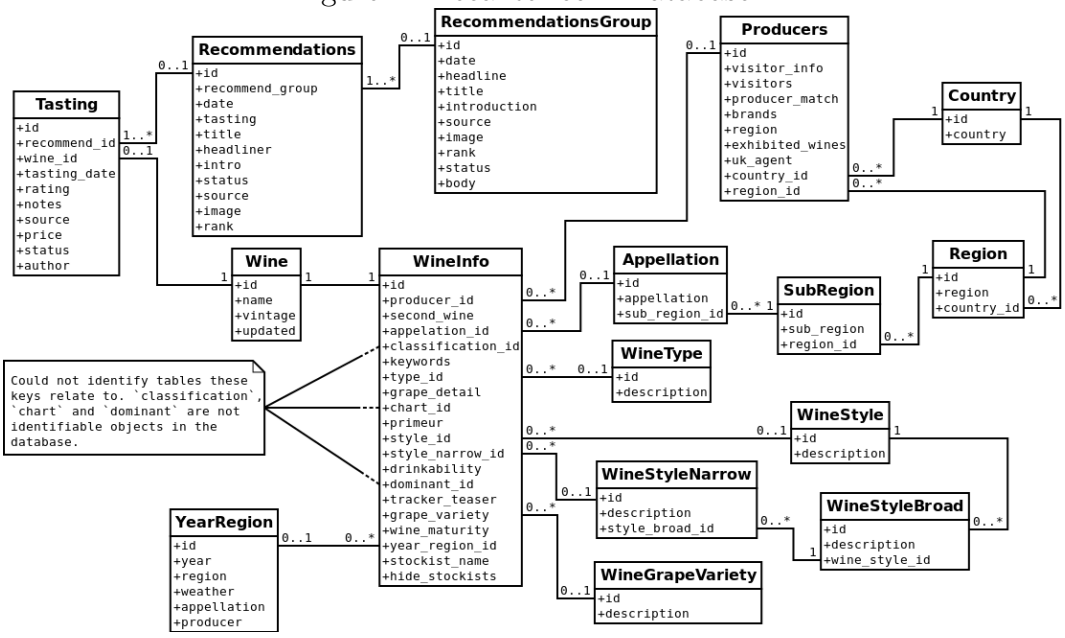

The "wine_info" table is a mixture of foreign keys joining to very small tables, such as "wine_info.type_id" joining to "wine_type.id" where "wine_type"

is a table with only two attributes. This approach, stiving for a high degree of normalisation, contrasts with the fact that the same table also has the attribute "second_wine", as string which only holds data in 450 of the 38762 entries in the table.

Creating The Sommelier Dataset

As well as containing a large amount of unreliable or incomplete data which was of little use to my project, the database's complex design would not lend itself to simple queries.

Rather than using a large number of foreign key relations, which would necessitate a large number of joins at query time, I hoped to pare down the database to the minimum number of tables for the maximum ease of retrieval and manipulation. Figure 2 shows the new database schema.

Figure 2: Sommelier Database

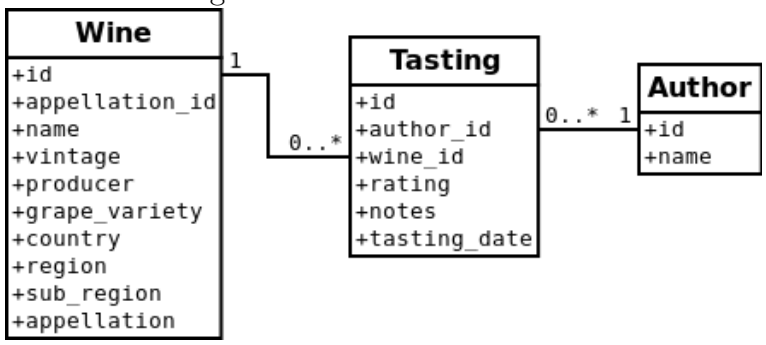

For the new Sommelier database, Figure 2, I decided to denormalise [explanation/citation needed!] the wine data. This enables the data to be queried without joins, maximising the simplicity and execution speed of the queries [citation needed]. Denormalization makes data integrity difficult to maintain however, as there are potentially a large number of records to update for any change in a duplicated value. In this case an appellation or sub-region name changing might require thousands of records to be amended. Creating and editing wines is not a requirement of my system however; for the purposes of this project the wine and tasting data is static. For this reason the duplication of data within the Wine records is not problematic, however in a real world setting this would need to be revisited.

The tables "wine_style_narrow" and "wine_style_broad" contained generic text descriptions for wine ("rich and creamy", "crisp and tangy" etc.). I initially considered this to have potential for migration into tag data which I could reuse as part of my filtering. Unfortunately less than 6435 of the records in "wine_info" had non-null values for their "style_narrow_id" field, and only 3397 had associations in the "tasting" table. 3397 was only around 10% of the wines I expected to be working with in my migrated dataset, I reckoned this too low a proportion to be worthwhile.

The table "wine_style" was rejected for the same reasons as those above.

The "wine_type" table was ignored because no wines were related to it. Those wines with the "type_id" field populated in the database were populated with keys which did not match any "id" in the type table.

The "producers" table was not disregarded, as there were over 20,000 wines with tasting notes and a producer record.

The Author Problem

The biggest shortcoming of the dataset is that the author of a tasting note is often not recorded. The number of wines with notes and known authors is only 1401, with there being 18 named authors on the system.

Table 1 shows the distribution of tastings amongst authors, only 5 of which have tasted and rated more than 100 wines in the database.

In some cases an author's initials or full name are recorded within the text of a tasting note. I decided that extracting and making use of these was impractical given the time constraints of this project.

DESCRIBE DATA SETS BEFORE AND AFTER

THE SOMMELIER DATASET

Having analysed the dataset and conceived an ideal schema, I needed to decide what the criteria to apply when extracting my new dataset from the source data.

Given that the purpose of the dataset is social recommendations, the first decision I made was to discard any wines without both tasting notes and a rating, whether.

Table 1: Authors of tasting notes and ratings

Author	Wines tasted, with notes and rating
Amy Wislocki	28
Andrew Jefford	105
Beverley Blanning MW	13
Carolyn Holmes	1
Christelle Guibert	119
Clive Coates MW	6
David Peppercorn	44
Gerald D Boyd	7
Harriet Waugh	250
James Lawther MW	226
John Radford	2
Josephine Butchart	24
Norm Roby	4
Rosemary George MW	6
Serena Sutcliffe	31
Stephen Brook	19
Steven Spurrier	497

5 Testing and Evaluation

How well does the system work? Details of testing and evaluation of the system...

6 Conclusion

Was the project successful?

7 Review

Review / reflections of the project on a personal level. What has been achieved? What were the problems, and how were they overcome?

Lessons learnt... - Data cleanup very time consuming - Literature vast -¿ plural techniques for recommendation: very difficult to work out what strategy is best for given situation.

References

[1] Burke, R., *The Wasabi Personal Shopper: A Case-Based Recommender System*, 1999. Submitted to the 11th Annual Conference on Innovative Applications of Artificial Intelligence.

[2] Burke, R., *Integrating Knowledge-Based and Collaborative-Filtering Recommender Systems*, 1999. In: Artificial Intelligence for Electronic Commerce: Papers from the AAAI Workshop (AAAI Technical Report WS-99-0 1), pp.69-72.

[3] Burke, R., *Knowledge-Based Recommender Systems*, Encyclopedia of Library and Information Systems, 2000. Marcel Dekker.

[4] Burke, R., *Hybrid Recommender Systems: Survey and Experiments*, User Modeling and User-Adapted Interaction, Volume 12 Issue 4, November 2002, Pages 331 - 370. Kluwer Academic Publishers: Hingham, MA, USA

[5] Debnath, Souvik and Ganguly, Niloy and Mitra, Pabitra, *Feature weighting in content based recommendation system using social network analysis*, Proceedings of the 17th international conference on World Wide Web, WWW '08, 2008, Beijing, China, Pages 1041 - 1042. ACM: New York, NY, USA,

[6] Goldberg, D. Nichols, D., Oki, B. M., and Terry, D., *Using collaborative filtering to weave an information tapestry*, Commun. ACM 35, 12 (Dec. 1992), 61–70.

[7] Mangalindan, J. P., *Amazon's Recommendation Secret*, July 2012. URL: http://tech.fortune.cnn.com/2012/07/30/amazon-5/

[8] Patton, E., McGuinness, D., *Scaling the Wall: Experiences Adapting a Semantic Web Application to Utilize Social Networks on Mobile Devices*, 2010. In: Proceedings of the WebSci10: Extending the Frontiers of Society On-Line, April 26-27th, 2010, Raleigh, NC: US.

[9] Resnick, P., Iacovou, N., Sushak, M., Bergstrom, P., Riedl, J., *GroupLens: An open architecture for collaborative filtering of netnews*, 1994 ACM Conference on Computer Supported Collaborative Work, 1994. Association of Computing Machinery, Chapel Hill, NC.

[10] Resnick, P., Varian, H. R., *Recommender Systems*, 1997. Communications of the ACM, 40 (3), 56-58. Association of Computing Machinery, Chapel Hill, NC.

[11] http://wineagent.tw.rpi.edu/index.php